

# Intelligent Agents




# Outline

- Agent function and agent program
- Rationality
- PEAS (Performance measure, Environment, Actuators, Sensors)
- Environment types
- Agent types

# Agents

- An **agent** is anything that can be viewed as **perceiving** its **environment** through **sensors** and **acting** upon that environment through **actuators**




# sensors/percepts and actuators/actions?

- Human agent:
  - eyes, ears, and other organs for sensors; hands, legs, mouth, and other body parts for actuators
- Robotic agent:
  - cameras and infrared range finders for sensors; various motors for actuators
- Software Agent:
  - keystrokes, file contents & network packets as sensory inputs; acts by displaying on screen, writing to files & sending network packets


# Terminologies

- **Percept**: the agent's perceptual inputs
- **Percept sequence**: the complete history of everything the agent has perceived
- **Agent function**: maps any given percept sequence to an action  $[f: p^* \rightarrow A]$
- **The agent program**: runs on the physical architecture to produce  $f$
- **Agent = architecture + program**



# Example: Vacuum-cleaner world

- **Percepts:**  
Location and status,  
e.g., [A, Dirty]
- **Actions:**  
Left, Right, Suck, NoOp




Example vacuum agent program:

```
function Vacuum-Agent([location,status]) returns an action  
  if status = Dirty then return Suck  
  else if location = A then return Right  
  else if location = B then return Left
```

# A Simple Agent Function

Percept sequence	Action
[A, Clean]	?
[A, Dirty]	
[B, Clean]	
[B, Dirty]	
[A, Clean], [A, Clean]	
[A, Clean], [A, Dirty]	
...	
[A, Clean], [A, Clean], [A, Clean]	
[A, Clean], [A, Clean], [A, Dirty]	
...	




# Rational agents

- "do the right thing"
- Select actions expected to maximize its **performance measure**
- Based on evidence provided by the percept sequence and the agent's built-in knowledge
- **Performance measure (utility function):**  
An *objective* criterion for success of an agent's behavior
- Definition:
  - For each possible **percept sequence**, a rational agent should select an **action** that is expected to maximize its **performance measure**, given the evidence provided by the percept sequence and whatever **built-in knowledge** the agent has.



# Example: Vacuum-Agent


- **Percepts:**  
Location and status,  
e.g., [A,Dirty]
- **Actions:**  
Left, Right, Suck, Dump, NoOp



- Potential performance measures for our vacuum agent?
  - amount of dirt cleaned up,
  - amount of time taken,
  - amount of electricity consumed,
  - amount of noise generated, etc.

# Example: Vacuum-Agent

- **Percepts:**  
Location and status,  
e.g., [A,Dirty]
- **Actions:**  
Left, Right, Suck, Dump, NoOp



**function Vacuum-Agent([location,status])** returns an **action**

- *if status = Dirty then return Suck*
- *else if location = A then return Right*
- *else if location = B then return Left*

Is this agent rational?

# Autonomy

- **Autonomous agent** – behavior determined by its own experience
  - ability to learn and adapt
  - can always say “no”
  - needs enough built-in knowledge to survive
- **Not autonomous agent**
  - guided by its designer

# Task Environment Specification

- Problem specification: **Performance measure, Environment, Actuators, Sensors (PEAS)**
- **Example: autonomous taxi**
  - **Performance measure**
    - Safe, fast, legal, comfortable trip, maximize profits
  - **Environment**
    - Roads, other traffic, pedestrians, customers
  - **Actuators**
    - Steering wheel, accelerator, brake, signal, horn
  - **Sensors**
    - Cameras, speedometer, GPS, odometer, engine sensors, keyboard

# Agent: Spam filter

- Performance measure
  - Minimizing false positives, false negatives
- Environment
  - A user's email account, email server
- Actuators
  - Mark as spam, delete, etc.
- Sensors
  - Incoming messages, other information about user's account

# Environment types


- Fully observable vs. partially observable
- Deterministic vs. stochastic
- Episodic vs. sequential
- Static vs. dynamic
- Discrete vs. continuous
- Single agent vs. multi-agent

# Fully observable vs. partially observable

- Do the agent's sensors give it access to the complete state of the environment?



VS.



# Deterministic vs. stochastic

- Is the next state of the environment completely determined by the current state and the agent's action?



VS.






# Episodic vs. sequential

- Is the agent's experience divided into unconnected episodes, or is it a coherent sequence of observations and actions?



VS.



# Static vs. dynamic

- Is the world changing while the agent is thinking?



VS.




# Discrete vs. continuous

- Does the environment provide a fixed number of distinct percepts, actions, and environment states?
  - Time can also evolve in a discrete or continuous fashion



vs.




# Single-agent vs. multiagent


- Is an agent operating by itself in the environment?



vs.



# Examples of different environments



# Word jumble solver



# Chess with a clock



# Scrabble



# Autonomous driving

Observable	Fully	Fully	Partially	Partially
Deterministic	Deterministic	Deterministic	Stochastic	Stochastic
Episodic	Episodic	Sequential	Sequential	Sequential
Static	Static	Dynamic	Static	Dynamic
Discrete	Discrete	Continuous	Discrete	Continuous
Single agent	Single	Multi	Multi	Multi

# Hierarchy of agent types

## (0) Table-driven agents

- percept sequence/action table to find the next action
- (large) **lookup table**

## (1) Simple reflex agents

- **condition-action rules**,
- production system; no memory of past world states

## (2) Agents with memory

- **internal state** to keep track of past states of the world

## (3) Agents with goals

- state and **goal information** about desirable situations
- take future events into consideration

## (4) Utility-based agents

- base decisions on **utility theory** in order to act rationally

simple



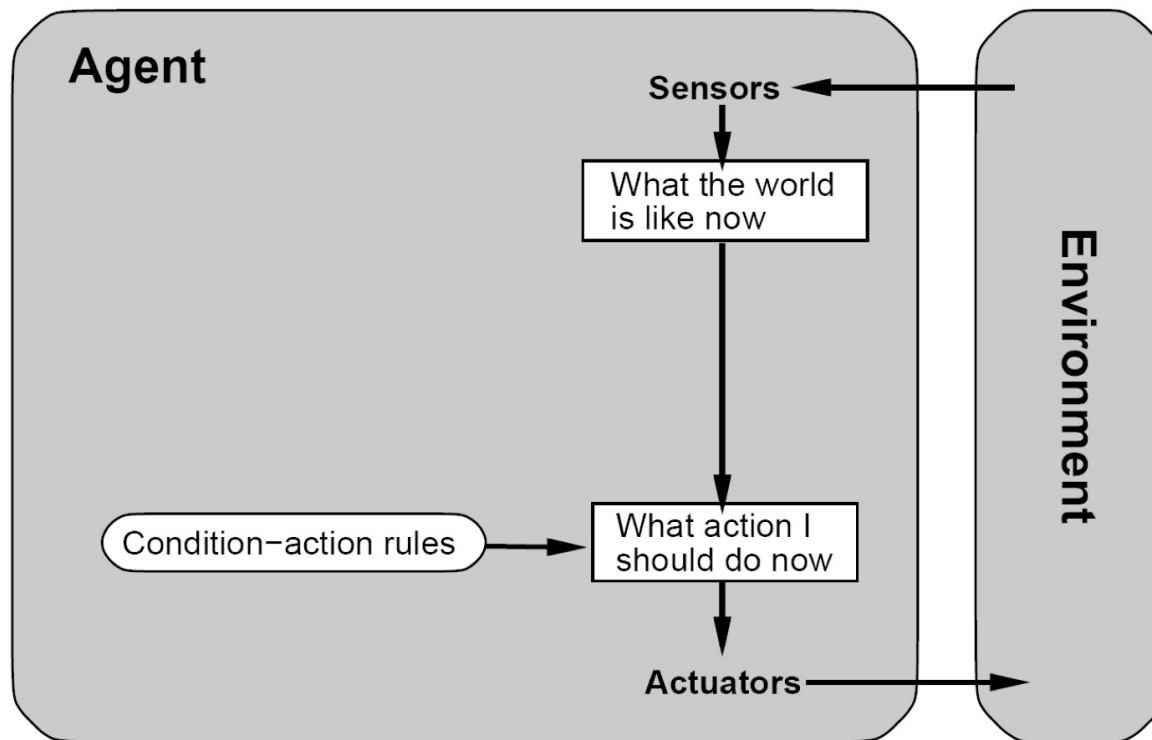
complex

# Table-Driven Agents - Problems

- Daunting sizes of tables
  - physical agents cannot store
  - designer – no time and knowledge to create
  - Not very practical!

# Simple reflex agent

- Select action based on current percept only
- Implemented through condition-action rules: if dirty then suck




Examples of classic **if-then** algorithms:

**if** (I sense a certain input) **then** (I apply a specific rule)

Issue: environment needs to be fully-observable. (Consider a self-driving car?)



# The vacuum-cleaner world



function REFLEX-VACUUM-AGENT (percept) returns an action

**static:** *rules*, a set of condition-action rules

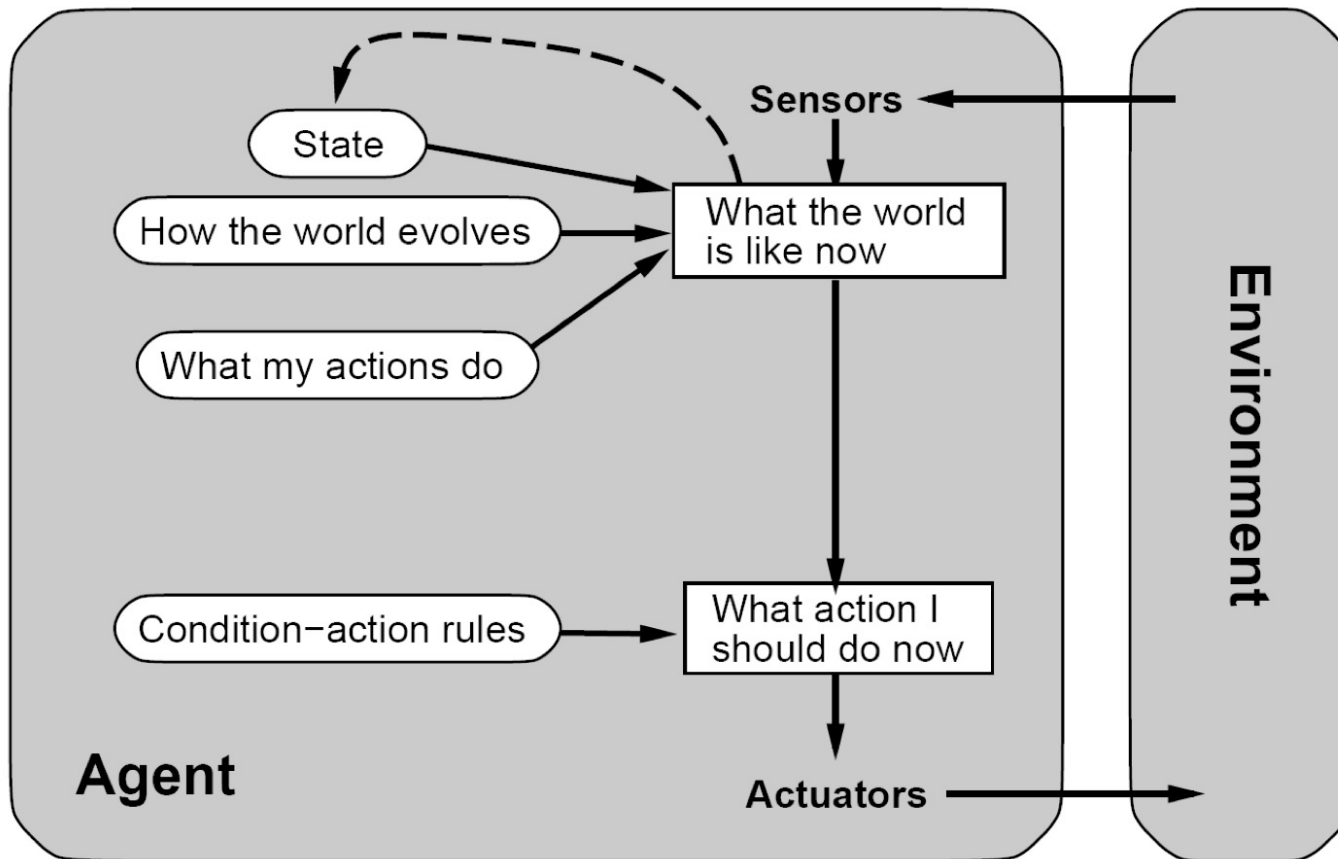
*rule*  $\leftarrow$  RULE-MATCH(*state*, *rules*)

*action*  $\leftarrow$  RULE-ACTION[*rule*]

return *action*

# Model-based reflex agent

- Internal state: aspects of the environment that cannot be currently observed
- For partially observable environments



# Model-based reflex agents

**function** REFLEX-AGENT-WITH-STATE(*percept*) **returns** an action

**static:** *rules*, a set of condition-action rules

*model*, a description of how the next state depends on current state and action

*state*, a description of the current world state

*action*, the most recent action

*state*  $\leftarrow$  UPDATE-STATE(*state*, *action*, *percept*, *model*)

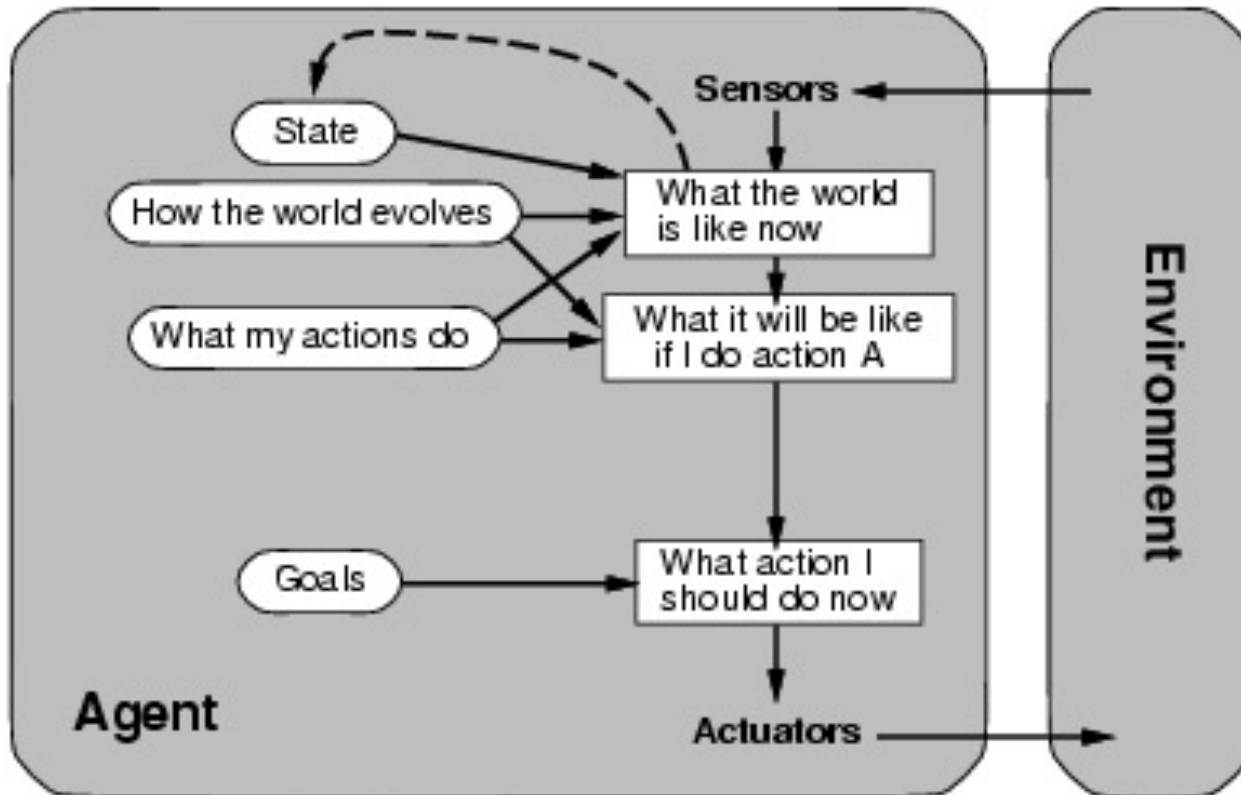
*rule*  $\leftarrow$  RULE-MATCH(*state*, *rules*)

*action*  $\leftarrow$  RULE-ACTION[*rule*]

**return** *action*

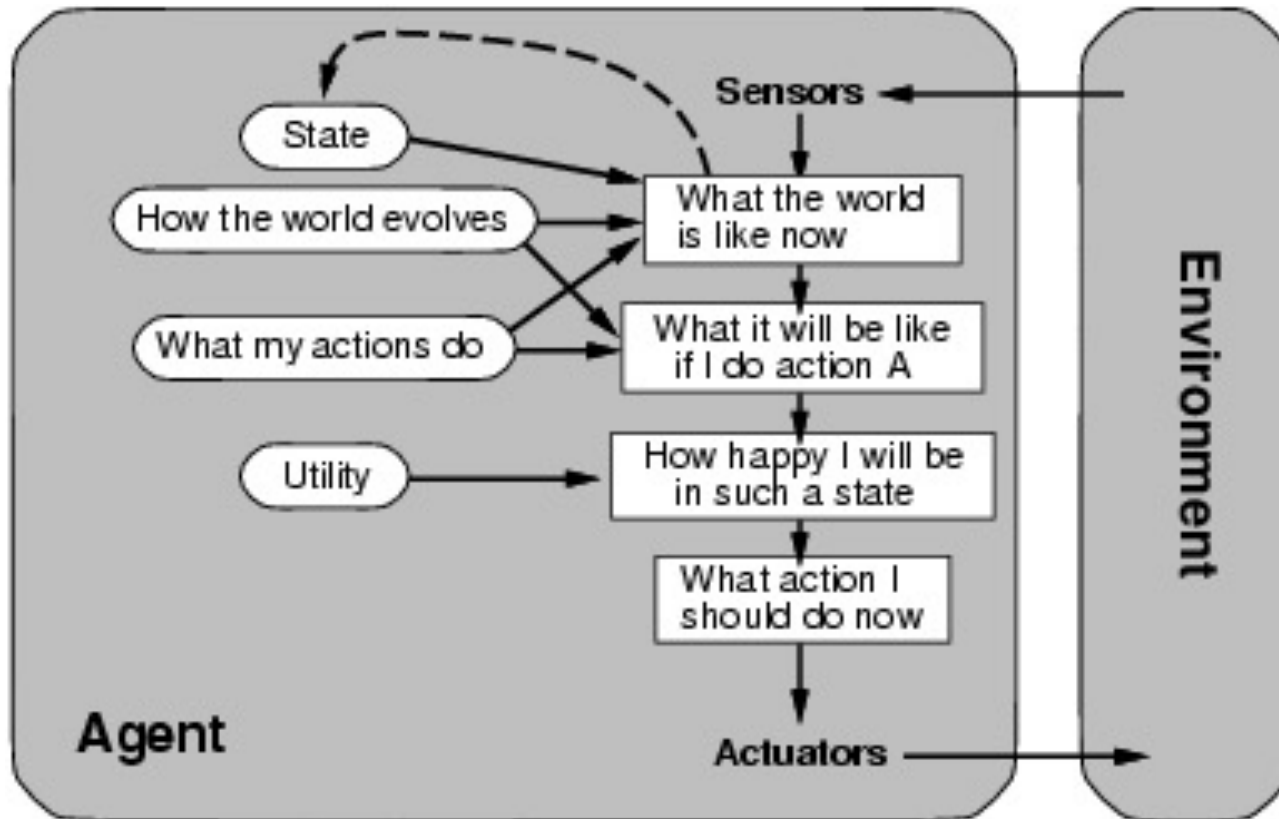
# Goal-based agent

- Goal information to select between possible actions in the current state
- Goal: desirable situations



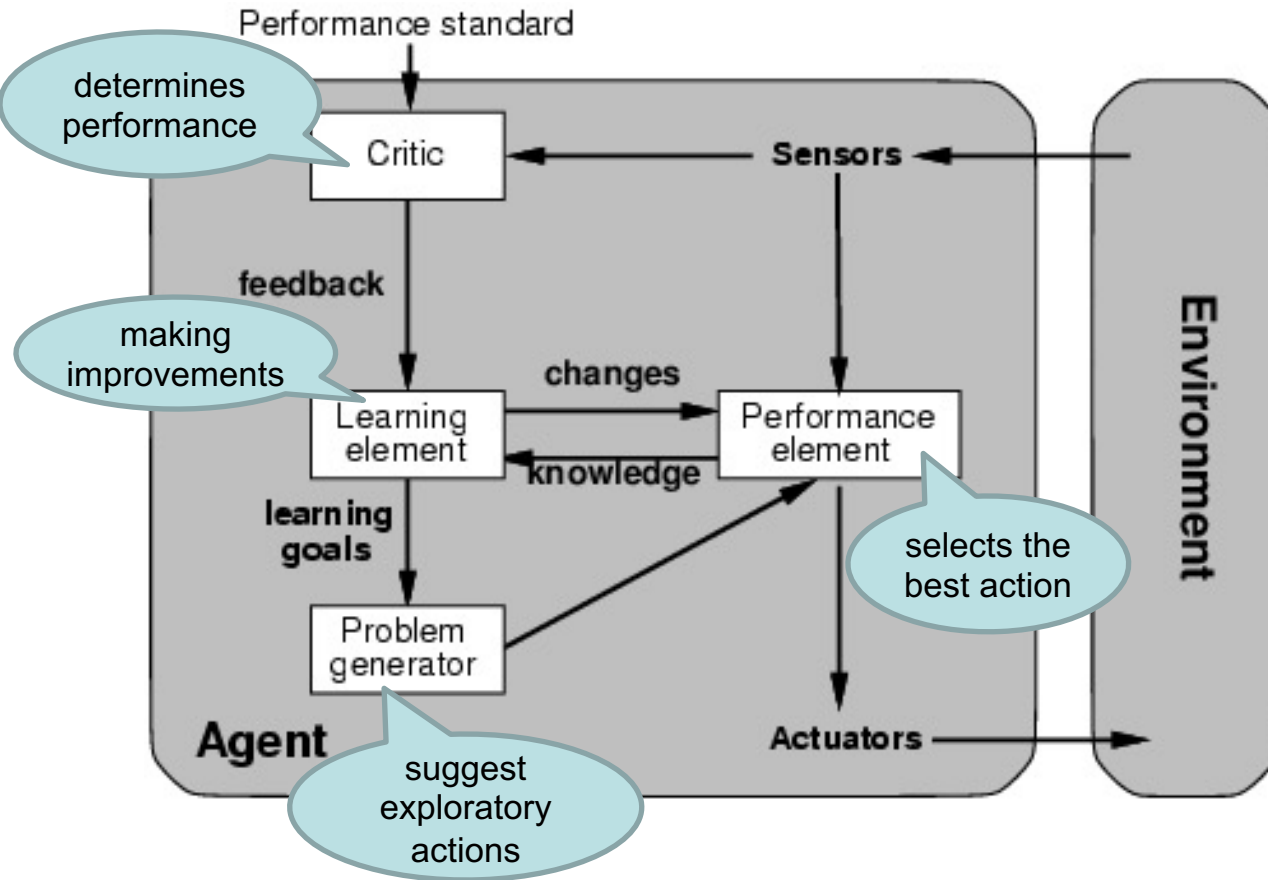
# Utility-based agent

- Utility function: evaluate desirability of states resulting from each possible action



- Certain goals can be reached in different ways:
  - “better” ways have a higher utilities
- Utility function: maps (sequence of) state(s) onto a real number.

# Learning/Autonomous agent



All previous agent-  
programs: methods for  
selecting *actions*

- Learning mechanisms can perform this task
- Teach instead of instructing
- Advantage: robustness in initially unknown environments.

# Summary

- An **agent** perceives and acts in an environment, has an architecture, and is implemented by an agent program
- A **rational agent** always chooses the action that maximizes its expected performance, given its percept sequence so far
- An **autonomous agent** uses its own experience rather than built-in knowledge of the environment by the designer
- An **agent program** maps percepts to actions and updates its internal state
  - **Reflex agents** respond immediately to percepts
  - **Goal-based agents** act in order to achieve their goal(s)
  - **Utility-based agents** maximize their own utility function
- The most challenging environments are **partially observable**, **stochastic**, **sequential**, **dynamic**, and **continuous**, and contain multiple intelligent agents.